

201709 evaluate C...

```
%pyspark
```

FINISHED

```
# Simple script to demonstrate evaluation of CommonCrawl-derived domain vectors by using
# classify domains according to high-level topic in the DMOZ dataset. Currently configured
# Bill's domain hex feature vectors from the 'Bill 6' notebook.
# TODO: Should we really be trying to predict domain links instead?
# PJ - 14 Sept 2017
```

```
import csv
import boto
from pyspark.sql.types import *
```

```
# Import the DMOZ domain category dataset
# (downloaded from https://dataverse.harvard.edu/dataset.xhtml?persistentId=doi:10.7910
```

```
dmoz_labels=sc.textFile('s3://billsdata.net/CommonCrawl/DMOZ/dmoz_domain_category.csv')
header = dmoz_labels.first() # extract header
dmoz_labels = dmoz_labels.filter(lambda row: row != header) # remove header row
dmoz_labels.take(1)
```

```
[u'"sdcastroverde.com"', "Top/World/Galego/regional/Galicia/Lugo/municipalities/Castroverde"]
```

```
%pyspark
```

FINISHED

```
# For now, collect all labels into a list on one node
# TODO: Could probably do this much faster using map-reduce!
dmoz_labels_list = dmoz_labels.collect()
```

```
# Take a look at one record
dmoz_labels_list[1]
```

```
u'"www.232analyzer.com"', "Top/Computers/Hardware/Test_Equipment/Analyzers"]
```

```
%pyspark
```

FINISHED

```
# Make a dictionary of short domains (without www.) to top-level category label, as per
# http://dmoztools.net
labels={}
prefix="www."
```

```
# TODO: Could probably do this much faster using map-reduce!
print(len(dmoz_labels_list))
for row in dmoz_labels_list[1:900000]: # Sample initially for speed (increasing to 1M c
```

```

row = row.replace(' ','').split(',')
fulldomain = row[0]
shortdomain = fulldomain[len(prefix):] if fulldomain.startswith(prefix) else fulldo
label = row[1].split("/")[-1].split("|")[0]
labels[shortdomain]=label
#print(shortdomain + " " + label)

```

```
# Take a look at the category for one domain from our dictionary
```

```

2488259
u'Computers'

```

```
%pyspark
```

FINISHED

```

# Summarize categories in the DMOZ data
from collections import Counter
Counter(labels.values())

```

```

Counter({u'World': 462130, u'Regional': 231978, u'Business': 54054, u'Society': 28807, u
'Arts': 24093, u'Shopping': 19102, u'Recreation': 16737, u'Computers': 16513, u'Sports':
12248, u'Science': 9857, u'Health': 8874, u'Reference': 7934, u'Games': 3730, u'Home': 2
549, u'News': 1383})

```

```
%pyspark
```

FINISHED

```

# Load Bill's domain feature vectors from s3, in the following format:
# (u'www.angelinajolin.com', [4.30406509320417, 0.02702702702702703, 0.0, 0.13513513513513513])

nfiles=128

# Load feature vectors from WAT files (from 'Bill 6' notebook):
inputURI = "s3://billsdata.net/CommonCrawl/domain_hex_feature_vectors_from_%d_WAT_files"
features_rdd = sc.textFile(inputURI).map(eval)
features_rdd.cache()
print("Nr domains:", features_rdd.count())
print(features_rdd.take(1))

```

FINISHED

[illegible]

FINISHED

Page 3 of 5

```

new_vec_ids=[]
new_vec_embs=[]
ground_truth=[]

def intersect(a, b):
    return list(set(a) & set(b))

common_domains=intersect(features_dict.keys(), labels.keys())
print(len(common_domains))
print(common_domains[1])

# Iterate over all the domain IDs for which we also have a vector embedding
for domain in common_domains:

    new_vec_ids.append(domain)
    new_vec_embs.append(features_dict[domain])
    ground_truth.append(labels[domain])

# Verify lengths of each list

3057
privateerpress.com
3057 3057 3057

```

%pyspark

FINISHED

```

# Split into training and test sets
from sklearn.cross_validation import train_test_split
X_train, X_test, y_train, y_test = train_test_split(new_vec_embs, ground_truth, test_si:

```

%pyspark

FINISHED

```

# Summarize labels in our test data
Counter(y_test)

Counter({u'World': 587, u'Regional': 283, u'Computers': 103, u'Arts': 87, u'Society': 86
, u'Business': 71, u'Reference': 69, u'Recreation': 49, u'Science': 41, u'Shopping': 29,
u'Sports': 28, u'Health': 25, u'Games': 25, u'News': 23, u'Home': 23})

```

%pyspark

FINISHED

```

# Fit KNN classifier to the training data and report results on test set
from sklearn.neighbors import KNeighborsClassifier
neigh = KNeighborsClassifier(n_neighbors=3, metric='cosine', algorithm='brute')
neigh.fit(X_train, y_train)
from sklearn.metrics import classification_report
print(classification_report(y_test, neigh.predict(X_test)))

```

	precision	recall	f1-score	support
Arts	0.05	0.13	0.07	87
Business	0.09	0.17	0.12	71
Computers	0.11	0.17	0.13	103
Games	0.02	0.04	0.03	25
Health	0.00	0.00	0.00	25
Home	0.03	0.04	0.04	23
News	0.06	0.04	0.05	23
Recreation	0.03	0.02	0.02	49
Reference	0.07	0.09	0.08	69
Regional	0.23	0.22	0.23	283
Science	0.11	0.05	0.07	41
Shopping	0.00	0.00	0.00	29
Society	0.16	0.06	0.09	86
Sports	0.00	0.00	0.00	28
World	0.50	0.38	0.43	587
avg / total	0.27	0.22	0.24	1529

`%pyspark`

FINISHED

```
# Fit Random Forest classifier to the training data and report results on test set
from sklearn.ensemble import RandomForestClassifier
rf = RandomForestClassifier(max_depth=2, random_state=0)
rf.fit(X_train, y_train)
print(classification_report(y_test, rf.predict(X_test)))
```

	precision	recall	f1-score	support
Arts	0.00	0.00	0.00	87
Business	0.00	0.00	0.00	71
Computers	0.00	0.00	0.00	103
Games	0.00	0.00	0.00	25
Health	0.00	0.00	0.00	25
Home	0.00	0.00	0.00	23
News	0.00	0.00	0.00	23
Recreation	0.00	0.00	0.00	49
Reference	0.00	0.00	0.00	69
Regional	0.00	0.00	0.00	283
Science	0.00	0.00	0.00	41
Shopping	0.00	0.00	0.00	29
Society	0.00	0.00	0.00	86
Sports	0.00	0.00	0.00	28
World	0.38	1.00	0.55	587
avg / total	0.15	0.38	0.21	1529

`%pyspark`

READY