

201709 evaluate C...

%pyspark

FINISHED

```
# Zeppelin notebook to demonstrate evaluation of CommonCrawl-derived domain vectors by
# classify domains according to high-level topic in the DMOZ dataset. Currently configu
# Bill's domain hex feature vectors from the 'Bill 6' notebook, and to use only Pyspark
# All cells should complete in less than a few minutes on an m4.2xlarge cluster.
# End-to-end run-time: approx 30 mins, with nfiles=128.
# NOTE: Should we really be trying to predict domain links instead? Or predicting bad d
# PJ - 20 Sept 2017
```

```
import boto
from pyspark.sql.types import *
```

```
# Import the DMOZ domain category dataset as an RDD
# (downloaded from https://dataverse.harvard.edu/dataset.xhtml?persistentId=doi:10.7910.
```

```
dmoz_labels=sc.textFile('s3://billsdata.net/CommonCrawl/DMOZ/dmoz_domain_category.csv')
header = dmoz_labels.first() # extract header
dmoz_labels = dmoz_labels.filter(lambda row: row != header).map(lambda row: row.replace
dmoz_labels.take(3)
```

```
[[u'sdcastroverde.com', u'Top/World/Galego/regional/Galicia/Lugo/municipalities/Castrove
rde'], [u'www.232analyzer.com', u'Top/Computers/Hardware/Test_Equipment/Analyzers'], [u'
zschachwitz-tischtennis.de', u'Top/World/Deutsch/Sport/ball_Sports/table_tennis/Teams/Ge
rmany/Saxony']]
```

%pyspark

FINISHED

```
# Convert our labels RDD into a Spark DataFrame with a schema - neither column can be Ni
schema=StructType([StructField("domain", StringType(), False), StructField("categories"
dmoz_labels_df=spark.createDataFrame(dmoz_labels,schema)
dmoz_labels_df.printSchema()
print(dmoz_labels_df.count())
dmoz_labels_df.show(1)
dmoz_labels_df.cache()
```

root

```
-- domain: string (nullable = false)
-- categories: string (nullable = false)
```

2488259

```
+-----+-----+
|          domain|          categories|
+-----+-----+
|sdcastroverde.com|Top/World/Galego/...|
+-----+-----+
```

only showing top 1 row

DataFrame[domain: string, categories: string]

%pyspark

FINISHED

```
# Make a dictionary of short domains (removing www. prefix) to top-level category label
prefix="www."
dmoz_labels_clean=dmoz_labels_df.rdd.map(lambda row: ((row['domain'][len(prefix):] if row
                                                         row['categories'].split("/")[1].:

dmoz_labels_df.unpersist()
schema=StructType([StructField("domain", StringType(), False), StructField("category", :
dmoz_labels_clean_df=spark.createDataFrame(dmoz_labels_clean,schema)
dmoz_labels_clean_df.show(2)
dmoz_labels_clean_df.cache()
```

```
+-----+-----+
|          domain| category|
+-----+-----+
|sdcastroverde.com|    World|
| 232analyzer.com|Computers|
+-----+-----+
only showing top 2 rows
DataFrame[domain: string, category: string]
```

%pyspark

FINISHED

```
# Summarize categories in the DMOZ data
dmoz_labels_clean_df.groupBy('category').count().show()
```

```
| category| count|
+-----+-----+
|Recreation| 46095|
|    World|1273970|
|   Science| 28138|
|    Home| 6952|
|Computers| 45194|
|   Sports| 34890|
|   Health| 24218|
|   Society| 82079|
| Shopping| 54062|
|Reference| 21663|
|    Games| 10246|
|    Arts| 66721|
| Business| 148144|
| Regional| 642176|
|    News| 3711|
+-----+-----+
```

%pyspark

FINISHED

```
# Load Bill's domain feature vectors from s3, in the following format:
# (u'www.angelinajolin.com', [4.30406509320417, 0.02702702702702703, 0.0, 0.13513513513513513])
```

```
nfiles=128 # (takes about 5 mins for 128 files)
```

```
# Load feature vectors from WAT files (from 'Bill 6' notebook) as an RDD:
inputURI = "s3://billsdata.net/CommonCrawl/domain_hex_feature_vectors_from_%d_WAT_files"
features_rdd = sc.textFile(inputURI).map(eval)
import pyspark.sql.types as typ
schema=StructType([StructField("domain", StringType(), False), StructField("vector", ArrayType(DoubleType(), False), False)])
features_df=spark.createDataFrame(features_rdd,schema)
features_df.cache()
print("Nr domains:", features_df.count())
print(features_df.show(1))
features_df.printSchema()
```

```
('Nr domains:', 2626203)
```

```
+-----+-----+
|      domain|      vector|
```

```
+-----+-----+
|www.iggl.de|[3.63758615972638...|
```

```
+-----+-----+
```

```
only showing top 1 row
```

```
None
```

```
root
```

```
 |-- domain: string (nullable = false)
```

```
 |-- vector: array (nullable = true)
```

```
    |-- element: double (containsNull = false)
```

```
%pyspark
```

FINISHED

```
# Spark.ML classifiers require VectorUDF type, rather than Array, so we need to convert
from pyspark.ml.linalg import Vectors, VectorUDT
vectorize=udf(lambda vs: Vectors.dense(vs), VectorUDT())
features_df = features_df.withColumn("vec", vectorize(features_df['vector'])).drop('vector')
print(features_df.show(1))
features_df.printSchema()
```

```
+-----+-----+
|      domain|      vec|
```

```
+-----+-----+
|www.iggl.de|[3.63758615972638...|
```

```
+-----+-----+
```

```
only showing top 1 row
```

```
None
```

```
root
```

```
 |-- domain: string (nullable = false)
```

```
 |-- vec: vector (nullable = true)
```

```
%pyspark
```

FINISHED

```
# Filter embeddings for only those vectors that have entries in the DMOZ dictionary (i.e.
#common_domains_df= features_df.join(dmoz_labels_clean_df, features_df.domain == dmoz_labels_clean_df.domain)
common_domains_df=features_df.join(dmoz_labels_clean_df, ["domain"]) # doesn't create e
```

```
common_domains_df.cache()
features_df.unpersist()
dmoz_labels_clean_df.unpersist()
print("Number of labelled domains = " + str(common_domains_df.count()))
common_domains_df.show(3)
```

Number of labelled domains = 57479

```
+-----+-----+-----+
| domain|          vec|category|
+-----+-----+-----+
| 1by.by|[4.74493212836325...| World|
| 360.ch|[5.42934562895444...| World|
| 6311a.com|[3.46573590279972...| World|
+-----+-----+-----+
```

only showing top 3 rows

root

```
-- domain: string (nullable = false)
-- vec: vector (nullable = true)
-- category: string (nullable = false)
```

%pyspark

FINISHED

```
# Create numeric indexes for our classes
from pyspark.ml.feature import IndexToString, StringIndexer, VectorIndexer
labelIndexer = StringIndexer(inputCol="category", outputCol="indexedCategory").fit(comm
```

```
# Split into training and test sets using spark.ML API
domains_train, domains_test = common_domains_df.randomSplit([0.7,0.3],seed=42)
domains_test.groupBy('category').count().show()
```

```
| category|count|
+-----+-----+
|Recreation| 419|
| World| 6469|
| Science| 564|
| Home| 211|
| Computers| 1168|
| Sports| 400|
| Health| 233|
| Society| 1016|
| Shopping| 373|
| Reference| 857|
| Games| 209|
| Arts| 928|
| Business| 795|
| Regional| 3460|
| News| 180|
+-----+-----+
```

%pyspark

FINISHED

[illegible]

FINISHED

```
# Select (prediction, true label) and compute test error
from pyspark.ml.evaluation import MulticlassClassificationEvaluator
evaluator1 = MulticlassClassificationEvaluator(labelCol="indexedCategory", predictionCol="prediction")
evaluator2 = MulticlassClassificationEvaluator(labelCol="indexedCategory", predictionCol="prediction")

accuracy = evaluator1.evaluate(predictions)
f1=evaluator2.evaluate(predictions)

print("Accuracy=%g, F1=%g" % (accuracy, f1))
```

Accuracy=0.37432, F1=0.203905

%pyspark

FINISHED