

# Paul 5 - faster do...

```
%pyspark
```

FINISHED

```
# Zeppelin notebook to create domain summaries based on the May/Jun/Jul 2017 CommonCrawl
# as per description here: http://commoncrawl.org/2017/08/webgraph-2017-may-june-july/
# PJ - 6 October 2017
```

```
import boto
from pyspark.sql.types import *
```

```
LIMIT=1000000 # TODO - remove temporary limit to run full summaries!
```

```
# Import the PLD vertices list as a DataFrame
pld_schema=StructType([StructField("ID", StringType(), False), StructField("PLD", StringType(), False)])
pld_txt=sc.textFile("s3://commoncrawl/projects/hyperlinkgraph/cc-main-2017-may-jun-jul/pld.txt")
temp_pld = pld_txt.map(lambda k: k.split()) # By default, splits on whitespace, which is correct
pld_df=temp_pld.toDF(pld_schema).limit(LIMIT).repartition(16)
pld_df.show(3)
pld_df.cache()
# Should have 91M domains
#print(pld_df.count())
```

```
+---+-----+
| ID|          PLD|
+---+-----+
|  4|    aaa.aaaaaa|
| 20|laarp.disruptaging|
| 36|    abbott.it|
+---+-----+
only showing top 3 rows
DataFrame[ID: string, PLD: string]
```

```
%pyspark
```

FINISHED

```
# Next import the PLD edges as a DataFrame
pld_edges_schema=StructType([StructField("src", LongType(), False), StructField("dst", LongType(), False)])
pld_edges_txt=sc.textFile("s3://commoncrawl/projects/hyperlinkgraph/cc-main-2017-may-jun-jul/pld_edges.txt")
temp_edges_pld = pld_edges_txt.map(lambda k: map(int, k.split())) # By default, splits on whitespace, which is correct
pld_edges_df=temp_edges_pld.toDF(pld_edges_schema).limit(LIMIT).repartition(16) # TODO - remove temporary limit
pld_edges_df.show(3)
pld_edges_df.cache()
```

```

+---+-----+
|src|      dst|
+---+-----+
| 24| 9258593|
| 27| 416120|
| 27|17255607|
+---+-----+
only showing top 3 rows
DataFrame[src: bigint, dst: bigint]

```

```
%pyspark
```

FINISHED

```

# Load the host-level graph vertices in the same way
host_schema=StructType([StructField("hostid", StringType(), False), StructField("host",
host_txt=sc.textFile("s3://commoncrawl/projects/hyperlinkgraph/cc-main-2017-may-jun-jul
temp_host = host_txt.map(lambda k: k.split()) # By default, splits on whitespace, which
host_df=temp_host.toDF(host_schema).limit(LIMIT).repartition(16) # TODO - remove tempor
host_df.show(3)
host_df.cache()
# Should have 1.3B hosts
#print(host_df.count())

+-----+-----+
|hostid|      host|
+-----+-----+
|      4| aaa.aaa.aaa.aaal
|     20|aaa.ldap-primaryl
|     36|      abb.brandl
+-----+-----+
only showing top 3 rows
DataFrame[hostid: string, host: string]

```

```
%pyspark
```

FINISHED

```

# Load in all harmonic centrality and page-ranks, and join based on reverse domain name
# Format: #hc_pos #hc_val #pr_pos #pr_val #host_rev
#pr_schema=StructType([StructField("hc_pos", StringType(), False), StructField("hc_val"
StringType(), False)])
pr_txt=sc.textFile("s3://commoncrawl/projects/hyperlinkgraph/cc-main-2017-may-jun-jul/d
header=pr_txt.first()
pr_txt=pr_txt.filter(lambda x: x!=header)
temp_pr = pr_txt.map(lambda k: k.split()) # By default, splits on whitespace, which is
pr_df=temp_pr.toDF(header.split()).withColumnRenamed("#host_rev", "host_rev").limit(LIMI
pr_df.show(3)
pr_df.cache()
#print(pr_df.count()) # Should be 91M

```

```

+-----+-----+-----+-----+-----+
|#hc_pos| #hc_val|#pr_pos|          #pr_val|    host_rev|
+-----+-----+-----+-----+-----+
|      5|21234402|      5| 0.00706371663388533|    com.google|
|     23|17323658|     47|0.000801430611613196|    com.vimeo|
|     44|16646934|    185|0.000141243073946399|org.wikimedia|
+-----+-----+-----+-----+-----+

```

only showing top 3 rows

```
DataFrame[#hc_pos: string, #hc_val: string, #pr_pos: string, #pr_val: string, host_rev: string]
```

```
%pyspark
```

FINISHED

```

# Debug partitioning of our 4 big dataframes
sc.getConf().getAll() #.mkString("\n")
print(pld_df.rdd.getNumPartitions())
print(pld_edges_df.rdd.getNumPartitions())
print(host_df.rdd.getNumPartitions())
pr_df.rdd.getNumPartitions()

```

```

16
16
16
16

```

```
%pyspark #--packages graphframes:graphframes:0.5.0-spark2.1-s_2.11
```

FINISHED

```
# We now have everything we need in these four dataframes to create the summaries we ne
```

```

# This code can't handle the complete edge lists, and produces this exception:
# java.lang.IllegalArgumentException: Size exceeds Integer.MAX_VALUE
#out_degrees=dict(pld_edges_df.groupBy("src").count().collect())
#in_degrees=dict(pld_edges_df.groupBy("dst").count().collect())
#print(out_degrees['846558'])
#print(in_degrees['846558'])

```

```

# Instead, just create RDDs and use lookup()
out_degrees=pld_edges_df.groupBy("src").count()
in_degrees=pld_edges_df.groupBy("dst").count()
out_degrees.show(3)
in_degrees.show(3)
#print(out_degrees.rdd.lookup(846558))
#print(in_degrees.rdd.lookup(846558))

```

```

+-----+-----+
|  src|count|
+-----+-----+
|   29|   34|
| 2214|  127|
|13518|   22|
+-----+-----+
only showing top 3 rows
+-----+-----+
|      dst|count|
+-----+-----+
|       29|    8|
|27481043|    2|
|66742183|    2|
+-----+-----+
only showing top 3 rows

```

```
%pyspark
```

FINISHED

```

# Next, we'll construct a local dictionary from of all the PLDS (key is the PLD, value
# This is our truth-table of known PLDs that we'll use when counting hosts
# This code can't handle the full PLD list and produces this exception:
# Stack trace: ExitCodeException exitCode=52
pld_lookup_table=dict(pld_df.rdd.map(lambda x: (x['PLD'], x['ID']))).collect()
print(pld_lookup_table["aaa.aaa"])

# Instead, just create an RDD and use lookup()
#pld_lookup_table=pld_df.rdd.map(lambda x: (x['PLD'], x['ID']))
#print(pld_lookup_table.lookup("aaa.aaa"))

# Next, broadcast this map so it's available on all the slave nodes - this seems to bre
#pld_map=sc.broadcast(pld_lookup_table)

```

2

```
%pyspark
```

FINISHED

```

# Returns a Boolean to say whether PLD is a hostname in itself
def is_a_pld(hostname):
    if hostname in pld_lookup_table:
        #if pld_lookup_table.filter(lambda a: a == hostname).count()>0:
            return True
    else:
        return False

# Define a function to do the hostname->pld conversion, if the pld exists in our dictio
def convert_hostname(hostname):
    # Return hostname as-is, if this is already a PLD
    if hostname in pld_lookup_table:
        #if pld_lookup_table.filter(lambda a: a == hostname).count()>0:
            return hostname
    # Otherwise we're going to have to split it up and test the parts

```

```

try:
    parts=hostname.split('.')
    if (len(parts)>4 and is_a_pld('.'.join(parts[0:4]))):
        return '.'.join(parts[0:4])
    if (len(parts)>3 and is_a_pld('.'.join(parts[0:3]))):
        return '.'.join(parts[0:3])
    if (len(parts)>2 and is_a_pld('.'.join(parts[0:2]))):
        return '.'.join(parts[0:2])
    if (len(parts)>1 and is_a_pld('.'.join(parts[0:1]))):
        return '.'.join(parts[0:1])
    return "ERROR" # Couldn't find a corresponding PLD - this should never happen!
except:
    return "ERROR"

# Test
print(convert_hostname("aaa.aaa"))
print(is_a_pld("aaa.aaa"))

```

```

aaa.aaa
True

```

```
%pyspark
```

FINISHED

```

# Now count the number of hosts per PLD in a scalable way, and create another dictionary
# Takes 5mins for first 10M rows -> approx 8 hours for all 1.3B rows?
count_table=host_df2.drop('hostid').rdd.map(lambda x: (convert_hostname(x['host']),1)).collectAsMap()
bool_table=host_df2.drop('hostid').rdd.map(lambda x: (x['host'], is_a_pld(x['host']))).collectAsMap()
print(count_table['aaa.aaa'])
print(bool_table['aaa.aaa'])
print(count_table['ERROR']) # Should be zero once we've loaded all the PLDs!

# TODO: Fix error in collect()
# java.lang.IllegalArgumentException: Size exceeds Integer.MAX_VALUE

```

```

6
True
7180422

```

```
%pyspark
```

FINISHED

```

from pyspark.sql.types import IntegerType
from pyspark.sql.functions import udf, col, when, lit

# Define a UDF to perform column-based lookup
def translate(mapping):
    def translate_(col):
        if not mapping.get(col):
            return 0
        else:
            return mapping.get(col)
    return udf(translate_, IntegerType())

```

```
# And a similar function for the Boolean map
def translate_bool(mapping):
    def translate_bool_(col):
        if not mapping.get(col):
            return False
        else:
            return mapping.get(col)
    return udf(translate_bool_, BooleanType())

# Insert our count column back into the host summary dataframe, along with a boolean to
# While we're at it, let's add in the in and out-degrees too, and an indicator of wheth
# crawled_test=when(col("OutDegree")==0, lit(False)).otherwise(lit(True))
pld_df_joined=pld_df.withColumn('NumHosts', translate(count_table)("PLD"))\
    .withColumn('PLDisHost?', translate_bool(bool_table)("PLD"))\
    #.withColumn('InDegree', translate(in_degrees)("ID"))\
    #.withColumn('OutDegree', translate(out_degrees)("ID"))\
    #.withColumn('Crawled?', crawled_test)
pld_df_joined.sort("NumHosts", ascending=False).show(100)
nld_df_joined.cache()
1293592|ar.com.publicargr...| 1377| true|
1308619| ar.com.ruoffl | 1370| false|
1 83250| ar.com.a-e-a| 1363| false|
1244970| ar.com.lyros| 1356| false|
1818378| at.webnode| 1347| true|
1385848| asia.6eha| 1344| false|
1324057| ar.com.speis| 1337| false|
1293172| ar.com.psdil | 1323| false|
1276548| ar.com.ostiz| 1322| false|
1357313| ar.com.zonq| 1294| false|
1349896| ar.com.vmnet| 1292| false|
1192232| ar.com.fsgsa| 1248| false|
1 75471| am.schoolsite| 1243| true|
1 51928| ai.nl| 1241| true|
1186154| ar.com.fimct| 1237| false|
+-----+-----+-----+-----+
only showing top 100 rows
DataFrame[ID: string, PLD: string, NumHosts: int, PLDisHost?: boolean]
```

```
%pyspark
```

FINISHED

```
# Join with in-degree and out-degree dataframes
pld_df_joined2=pld_df_joined.join(out_degrees, out_degrees.src==pld_df_joined.ID, "left"
pld_df_joined.unpersist()
pld_df_joined3=pld_df_joined2.join(in_degrees, in_degrees.dst==pld_df_joined2.ID, "left"
pld_df_joined2.unpersist()
pld_df_joined3.show(5)
pld_df_joined3.cache()
```

```

+---+-----+-----+-----+-----+-----+
| ID|          PLD|NumHosts|PLDisHost?|OutDegree|InDegree|
+---+-----+-----+-----+-----+-----+
| 26|      abb.nic|      3|      true|      2|     null|
| 29|abbott.corelabora...|      2|      true|     34|       8|
| 474|    ac.americancars|      1|      true|     null|     null|
| 964|          ac.cmt|      1|     false|      1|     null|
|1677|    ac.insight|      1|      true|      7|     null|
+---+-----+-----+-----+-----+-----+

```

only showing top 5 rows

DataFrame[ID: string, PLD: string, NumHosts: int, PLDisHost?: boolean, OutDegree: bigint, InDegree: bigint]

```
%pyspark
```

FINISHED

```

# Insert a flag to indicate whether the PLD has been crawled
crawled_test=when(col("OutDegree").isNull(), lit(False)).otherwise(lit(True))
pld_df_joined4=pld_df_joined3.withColumn('Crawled?', crawled_test)
pld_df_joined3.unpersist()
pld_df_joined4.show(5)
pld_df_joined4.cache()

```

```

+---+-----+-----+-----+-----+-----+-----+
| ID|          PLD|NumHosts|PLDisHost?|OutDegree|InDegree|Crawled?|
+---+-----+-----+-----+-----+-----+-----+
| 26|      abb.nic|      3|      true|      2|     null|     true|
| 29|abbott.corelabora...|      2|      true|     34|       8|     true|
| 474|    ac.americancars|      1|      true|     null|     null|    false|
| 964|          ac.cmt|      1|     false|      1|     null|     true|
|1677|    ac.insight|      1|      true|      7|     null|     true|
+---+-----+-----+-----+-----+-----+-----+

```

only showing top 5 rows

DataFrame[ID: string, PLD: string, NumHosts: int, PLDisHost?: boolean, OutDegree: bigint, InDegree: bigint, Crawled?: boolean]

```
%pyspark
```

FINISHED

```

# Finally, join with the harmonic centrality and page-rank for each domain
# Note: could probably speed this up using something like above techniques, or by presorting
pld_df_joined5=pld_df_joined4.join(pr_df, pr_df.host_rev==pld_df_joined4.PLD, "leftOuter")
                                .withColumnRenamed("#hc_val", "HarmonicCentrality").withColumnRenamed("#pr_val", "PageRank")
pld_df_joined4.unpersist()
pld_df_joined5.show(5)
pld_df_joined5.cache()

```

```

+---+-----+-----+-----+-----+-----+-----+-----+
-----+-----+
| ID|          PLD|NumHosts|PLDisHost?|OutDegree|InDegree|Crawled?|HarmonicCentr
ality|PageRank|
+---+-----+-----+-----+-----+-----+-----+-----+
-----+-----+
| 120|          abc.web|      1|    false|    null|    null|    false|
null|    null|
| 311|          ac.8411|      1|    false|    null|      1|    false|
null|    null|
| 713|          ac.bgc|      1|    false|    null|    null|    false|
null|    null|
| 871|          ac.casinos|      1|     true|      2|    null|     true|
null|    null|
|1014|ac.cosmopolitanun...|      1|     true|    null|    null|    false|
null|    null|
+---+-----+-----+-----+-----+-----+-----+-----+
-----+-----+

```

```
%pyspark
```

FINISHED

```

# Save final table to S3 in compressed CSV format
outputURI="s3://billsdata.net/CommonCrawl/domain_summaries2/"
codec="org.apache.hadoop.io.compress.GzipCodec"
pld_df_joined5.coalesce(1).write.format('com.databricks.spark.csv').options(header='true',

```

```
%pyspark
```

READY