

# Paul 5 - faster do...

%pyspark

FINISHED

```
# Zeppelin notebook to create domain summaries based on the May/Jun/Jul 2017
  CommonCrawl graph
# as per description here: http://commoncrawl.org/2017/08/webgraph-2017-may-june-july/
# PJ - 4 October 2017
```

```
import boto
from pyspark.sql.types import *
```

```
LIMIT=1000000 # TODO - remove temporary limit to run full summaries!
```

```
# Import the PLD vertices list as a DataFrame
pld_schema=StructType([StructField("ID", StringType(), False), StructField("PLD",
  StringType(), False)])
pld_txt=sc.textFile("s3://commoncrawl/projects/hyperlinkgraph/cc-main-2017-may-jun-jul
  /domaingraph/vertices.txt.gz")
temp_pld = pld_txt.map(lambda k: k.split()) # By default, splits on whitespace, which
  is what we want
pld_df=temp_pld.toDF(pld_schema).limit(LIMIT)
```

```
+---+-----+
```

```
| ID|    PLD|
```

```
+---+-----+
```

```
|  0|  aaa.a|
```

```
|  1| aaa.aa|
```

```
|  2|aaa.aaa|
```

```
+---+-----+
```

```
only showing top 3 rows
```

```
DataFrame[ID: string, PLD: string]
```

%pyspark

FINISHED

```
# Next import the PLD edges as a DataFrame
pld_edges_schema=StructType([StructField("src", StringType(), False), StructField
  ("dst", StringType(), False)])
pld_edges_txt=sc.textFile("s3://commoncrawl/projects/hyperlinkgraph/cc-main-2017-may
  -jun-jul/domaingraph/edges.txt.gz")
temp_edges_pld = pld_edges_txt.map(lambda k: k.split()) # By default, splits on
  whitespace, which is what we want
```

```

+---+-----+
|src|      dst|
+---+-----+
| 2| 9193244|
|20|75600973|
|21|46356172|
+---+-----+
only showing top 3 rows
DataFrame[src: string, dst: string]

```

```
%pyspark
```

FINISHED

```

# Load the host-level graph vertices in the same way
host_schema=StructType([StructField("hostid", StringType(), False), StructField("host"
, StringType(), False)])
host_txt=sc.textFile("s3://commoncrawl/projects/hyperlinkgraph/cc-main-2017-may-jun
-jul/hostgraph/vertices.txt.gz")
temp_host = host_txt.map(lambda k: k.split()) # By default, splits on whitespace,
which is what we want
host_df=temp_host.toDF(host_schema).limit(LIMIT*10) # TODO - remove temporary limit to
run full summaries!

```

```

+-----+-----+
|hostid|  host|
+-----+-----+
|      0| aaa.a|
|      1| aaa.aal
|      2|aaa.aaal
+-----+-----+
only showing top 3 rows
DataFrame[hostid: string, host: string]

```

```
%pyspark
```

FINISHED

```

# Load in all harmonic centrality and page-ranks, and join based on reverse domain
name
# Format: #hc_pos #hc_val #pr_pos #pr_val #host_rev
#pr_schema=StructType([StructField("hc_pos", StringType(), False), StructField
("hc_val", StringType(), False), StructField("pr_pos", StringType(), False),
StructField("pr_val", StringType(), False), StructField("host_rev", StringType(),
False)])
pr_txt=sc.textFile("s3://commoncrawl/projects/hyperlinkgraph/cc-main-2017-may-jun-jul
/domaingraph/ranks.txt.gz")
header=pr_txt.first()
pr_txt=pr_txt.filter(lambda x: x!=header)
temp_pr = pr_txt.map(lambda k: k.split()) # By default, splits on whitespace, which is

```

```

+-----+-----+-----+-----+-----+
|#hc_pos| #hc_val|#pr_pos|          #pr_val|    host_rev|
+-----+-----+-----+-----+-----+
|      1|24989952|      1| 0.0155264576161686|  com.facebook|
|      2|22460880|      3|0.00866038900847366|  com.twitter|
|      3|22097514|      2| 0.0128827315785546|com.googleapis|
+-----+-----+-----+-----+-----+

```

only showing top 3 rows

```
DataFrame[#hc_pos: string, #hc_val: string, #pr_pos: string, #pr_val: string, host_rev:
string]
```

```
%pyspark #--packages graphframes:graphframes:0.5.0-spark2.1-s_2.11 RUNNING 0%
```

```
# We now have everything we need in these four dataframes to create the summaries we
  need.
```

```
# First, let's use the PLD edges dataframe to compute in and out-degrees for each PLD
  ID, and store as dictionaries.
```

```
# Note: we could use GraphFrames for this but it's a pain to get GraphFrames working
  in a Zeppelin notebook!
```

```
out_degrees=dict(pld_edges_df.groupBy("src").count().collect())
```

Started a minute ago.

```
%pyspark
```

PENDING

```
# Next, we'll construct a local dictionary from of all the PLDS (key is the PLD, value
  is the ID)
```

```
# This is our truth-table of known PLDs that we'll use when counting hosts
```

```
pld_lookup_table=dict(pld_df.rdd.map(lambda x: (x['PLD'], x['ID'])).collect())
print(pld_lookup_table["aaa.aaa"])
```

```
# Next, broadcast this map so it's available on all the slave nodes - this seems to
```

```
%pyspark
```

PENDING

```
# Define a function to do the hostname->pld conversion, if the pld exists in our
  dictionary
```

```
def convert_hostname(hostname):
```

```
    # Return hostname as-is, if this is already a PLD
```

```
    if hostname in pld_lookup_table:
        return hostname
```

```
    # Otherwise we're going to have to split it up and test the parts
    try:
```

```
        parts=hostname.split('.')

```

```
        if (len(parts)>4 and '.'.join(parts[0:4]) in pld_lookup_table):
            return '.'.join(parts[0:4])

```

```
        if (len(parts)>3 and '.'.join(parts[0:3]) in pld_lookup_table):
            return '.'.join(parts[0:3])

```

```
        if (len(parts)>2 and '.'.join(parts[0:2]) in pld_lookup_table):
            return '.'.join(parts[0:2])

```

```

        if (len(parts)>1 and '.'.join(parts[0:1]) in pld_lookup_table):
            return '.'.join(parts[0:1])
        return "ERROR" # Couldn't find a corresponding PLD - this should never happen!
    except:
        return "ERROR"

```

# Returns a Boolean to say whether PLD is a hostname in itself

```

def is_hostname_a_pld(hostname):
    if hostname in pld_lookup_table:
        return True
    else:
        return False

```

# Test

```
print(convert_hostname("aaa.aaa"))
```

aaa.aaa

True

%pyspark

PENDING

# Now count the number of hosts per PLD in a scalable way, and create another dictionary

# Takes 5mins for first 10M rows -> approx 8 hours for all 1.3B rows?

```
count_table=host_df.drop('hostid').rdd.map(lambda x: (convert_hostname(x['host']),1)).reduceByKey(lambda x,y: x+y).collectAsMap()
```

```
bool_table=host_df.drop('hostid').rdd.map(lambda x: (x['host'], is_hostname_a_pld(x['host']))).filter(lambda x: x[1]==True).collectAsMap()
```

6

True

7180422

%pyspark

PENDING

```
from pyspark.sql.types import IntegerType
```

```
from pyspark.sql.functions import udf, col, when, lit
```

# Define a UDF to perform column-based lookup

```
def translate(mapping):
```

```
    def translate_(col):
```

```
        if not mapping.get(col):
```

```
            return 0
```

```
        else:
```

```
            return mapping.get(col)
```

```
    return udf(translate_, IntegerType())
```

# And a similar function for the Boolean map

```
def translate_bool(mapping):
```

```
    def translate_bool_(col):
```

```
        if not mapping.get(col):
```

```
            return False
```

```
        else:
```

```

    return mapping.get(col)
return udf(translate_bool_, BooleanType())

```

```

# Insert our count column back into the host summary dataframe, along with a boolean
  to say whether the PLD is a host in itself
# While we're at it, let's add in the in and out-degrees too, and an indicator of
  whether the site has been crawled.

```

```

crawled_test=when(col("OutDegree")==0, lit(False)).otherwise(lit(True))
pld_df_joined=pld_df.withColumn('NumHosts', translate(count_table)("PLD"))\
    .withColumn('PLDisHost?', translate_bool(bool_table)("PLD"))\
    .withColumn('InDegree', translate(in_degrees)("ID"))\
    .withColumn('OutDegree', translate(out_degrees)("ID"))\

```

```

1293592|ar.com.publicagr...| 1377| true| 1| 14| true|
1308619| ar.com.ruoffl| 1370| false| 0| 0| false|
1 83250| ar.com.a-e-al| 1363| false| 0| 0| false|
1244970| ar.com.lyros| 1356| false| 0| 0| false|
1818378| at.webnode| 1347| true| 372| 988| true|
1385848| asia.6eha| 1344| false| 0| 0| false|
1324057| ar.com.speis| 1337| false| 0| 0| false|
1293172| ar.com.psdil| 1323| false| 0| 0| false|
1276548| ar.com.ostiz| 1322| false| 0| 0| false|
1357313| ar.com.zonq| 1294| false| 0| 0| false|
1349896| ar.com.vmnet| 1292| false| 0| 0| false|
1192232| ar.com.fsgsa| 1248| false| 0| 0| false|
1 75471| am.schoolsit| 1243| true| 5| 95| true|
1 51928| ai.nl| 1241| true| 0| 0| false|
1186154| ar.com.fimct| 1237| false| 0| 0| false|

```

only showing top 100 rows

DataFrame[ID: string, PLD: string, NumHosts: int, PLDisHost?: boolean, InDegree: int, Ou

```
%pyspark
```

PENDING

```

# Finally, join with the harmonic centrality and page-rank for each domain
# Note: could probably speed this up using something like above techniques, or by
  presorting (but we don't really need to since this is only 91Mx91M)
pld_df_joined2=pld_df_joined.join(pr_df, pr_df.host_rev==pld_df_joined.PLID,
  "leftOuter").drop("#hc_pos").drop("#pr_pos").drop("host_rev").withColumnRenamed
  ("#hc_val", "HarmonicCentrality").withColumnRenamed("#pr_val", "PageRank")

```

110891	ac.d1bru1	11	true1	1201266614.49359706049864e-091
114351	ac.gorilla1	11	true1	911425614.50619088846452e-091
124761	ac.philter1	11	true1	1043478514.44625601709852e-091
131381	ac.ula1	11	false1	1004653114.5521807953781e-091
131451	ac.umedalen1	21	true1	1209300914.69402844088012e-091
133731	ac.yuil	21	true1	1210521715.09160908953513e-091
134841	academy.alphastar1	11	true1	981691914.91209890117111e-091
137681	academy.cirulnik1	11	true1	7967335.516.12220241367047e-091
137871	academy.cocoa1	11	true1	1019611916.02409896712952e-091
138821	academy.dental-coach1	11	true1	1098149514.43152551513855e-091
141571	academy.ger1	11	true1	808704711.37105166301547e-081
144101	academy.investmen...1	21	true1	1229918011.31584108265683e-081
147691	academy.newtown1	11	true1	1033442518.54845369040491e-091
152241	academy.talk1	21	true1	1201266715.30628234199031e-091
160641	accountant.buy-mo...1	11	true1	1019266115.29716249291459e-091

+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+  
only showing top 20 rows

DataFrame[ID: string, PLD: string, NumHosts: int, PLDisHost?: boolean, HarmonicCentralit

```
%pyspark
```

PENDING

```
# Save final table to S3 in compressed CSV format
outputURI="s3://billsdata.net/CommonCrawl/domain_summaries/"
codec="org.apache.hadoop.io.compress.GzipCodec"
pld_df_joined2.coalesce(1).write.format('com.databricks.spark.csv').options(header
```

```
%pyspark
```

FINISHED