

Paul 5 - faster do...

%pyspark

FINISHED

```
# Zeppelin notebook to create domain summaries based on the May/Jun/Jul 2017 CommonCrawl
# as per description here: http://commoncrawl.org/2017/08/webgraph-2017-may-june-july/
# PJ - 10 October 2017
```

```
import boto
from pyspark.sql.types import *
```

```
#LIMIT=10000000 # Temporary limit while developing code.
```

```
# Import the PLD vertices list as a DataFrame
#pld_schema=StructType([StructField("ID", StringType(), False), StructField("PLD", Strin
#pld_txt=sc.textFile("s3://commoncrawl/projects/hyperlinkgraph/cc-main-2017-may-jun-jul
#temp_pld = pld_txt.map(lambda k: k.split()) # By default, splits on whitespace, which
#pld_df=temp_pld.toDF(pld_schema) #.limit(LIMIT) #.repartition(4)
#pld_df.show(3)
```

```
# Load in an uncompressed, partitioned format, for fast reading in the future
saveURI="s3://billsdata.net/CommonCrawl/hyperlinkgraph/cc-main-2017-may-jun-jul/domaing
#pld_df.coalesce(64).write.save(saveURI) # Use all default options
pld_df=spark.read.load(saveURI)
pld_df.show(3)
pld_df.cache()
print(pld_df.count()) # Should have 91M domains
```

```
+---+-----+
| ID|    PLD|
+---+-----+
|  0|  aaa.a|
|  1|  aaa.a|
|  2|aaa.a|
+---+-----+
only showing top 3 rows
91034128
```

%pyspark

FINISHED

```
# Next import the PLD edges as a DataFrame
#pld_edges_schema=StructType([StructField("src", LongType(), False), StructField("dst",
#pld_edges_txt=sc.textFile("s3://commoncrawl/projects/hyperlinkgraph/cc-main-2017-may-j
#temp_edges_pld = pld_edges_txt.map(lambda k: map(int, k.split())) # By default, splits
#pld_edges_df=temp_edges_pld.toDF(pld_edges_schema) #.limit(LIMIT*10) #.repartition(8)
#pld_edges_df.show(3)
```

```
# Load in an uncompressed, partitioned format, for fast reading in the future
saveURI="s3://billsdata.net/CommonCrawl/hyperlinkgraph/cc-main-2017-may-jun-jul/domaing
#pld_edges_df.coalesce(64).write.save(saveURI) # Use all default options
```

```
pld_edges_df=spark.read.load(saveURI)
pld_edges_df.show(3)
pld_edges_df.cache()
```

```
+---+-----+
```

```
|src|      dst|
```

```
+---+-----+
```

```
| 2| 9193244|
```

```
|20|75600973|
```

```
|21|46356172|
```

```
+---+-----+
```

```
only showing top 3 rows
```

```
DataFrame[src: bigint, dst: bigint]
```

```
%pyspark
```

FINISHED

```
# Load the host-level graph vertices in the same way
#host_schema=StructType([StructField("hostid", StringType(), False), StructField("host"
#host_txt=sc.textFile("s3://commoncrawl/projects/hyperlinkgraph/cc-main-2017-may-jun-jul/
#temp_host = host_txt.map(lambda k: k.split()) # By default, splits on whitespace, which
#host_df=temp_host.toDF(host_schema) #.repartition(4)
#host_df.show(3)
```

```
# Save in an uncompressed, partitioned format, for fast reading in the future
saveURI="s3://billsdata.net/CommonCrawl/hyperlinkgraph/cc-main-2017-may-jun-jul/hostgra
#host_df.coalesce(128).write.save(saveURI) # Use all default options
host_df=spark.read.load(saveURI)
host_df.show(3)
host_df.cache()
#print(host_df.count()) # Should have 1.3B hosts
```

```
+-----+-----+
```

```
|hostid|  host|
```

```
+-----+-----+
```

```
|      0| aaa.a|
```

```
|      1| aaa.a|
```

```
|      2|aaa.a|
```

```
+-----+-----+
```

```
only showing top 3 rows
```

```
DataFrame[hostid: string, host: string]
```

```
%pyspark
```

FINISHED

```
# Load in all harmonic centrality and page-ranks, and join based on reverse domain name
# Format: #hc_pos #hc_val #pr_pos #pr_val #host_rev
#pr_schema=StructType([StructField("hc_pos", StringType(), False), StructField("hc_val"
(), False), StructField("host_rev", StringType(), False)])
#pr_txt=sc.textFile("s3://commoncrawl/projects/hyperlinkgraph/cc-main-2017-may-jun-jul/
#header=pr_txt.first()
#pr_txt=pr_txt.filter(lambda x: x!=header)
#temp_pr = pr_txt.map(lambda k: k.split()) # By default, splits on whitespace, which is
#pr_df=temp_pr.toDF(header.split()).withColumnRenamed("#host_rev", "host_rev") #.limit(L
```

```
#pr_df.show(3)
```

```
# Save in an uncompressed, partitioned format, for fast reading in the future
saveURI="s3://billsdata.net/CommonCrawl/hyperlinkgraph/cc-main-2017-may-jun-jul/domaining
#pr_df.coalesce(64).write.save(saveURI) # Use all default options
pr_df=spark.read.load(saveURI)
pr_df.show(3)
pr_df.cache()
```

```
+-----+-----+-----+-----+-----+
|#hc_pos| #hc_val|#pr_pos|          #pr_val|    host_rev|
+-----+-----+-----+-----+-----+
|      1|24989952|      1| 0.0155264576161686|  com.facebook|
|      2|22460880|      3| 0.00866038900847366|  com.twitter|
|      3|22097514|      2| 0.0128827315785546| com.googleapis|
+-----+-----+-----+-----+-----+
```

only showing top 3 rows

```
DataFrame[#hc_pos: string, #hc_val: string, #pr_pos: string, #pr_val: string, host_rev:
string]
```

```
%pyspark
```

FINISHED

```
# Debug partitioning of our 4 big dataframes
sc.getConf().getAll() #.mkString("\n")
print(pld_df.rdd.getNumPartitions())
print(pld_edges_df.rdd.getNumPartitions())
print(host_df.rdd.getNumPartitions())
pr_df.rdd.getNumPartitions()
```

```
64
64
117
64
```

```
%pyspark #--packages graphframes:graphframes:0.5.0-spark2.1-s_2.11
```

FINISHED

```
# We now have everything we need in these four dataframes to create the summaries we ne
```

```
# This code can't handle the complete edge lists, and produces this exception:
# java.lang.IllegalArgumentException: Size exceeds Integer.MAX_VALUE
#out_degrees=dict(pld_edges_df.groupBy("src").count().collect())
#in_degrees=dict(pld_edges_df.groupBy("dst").count().collect())
#print(out_degrees['846558'])
#print(in_degrees['846558'])
```

```
# Instead, just create RDDs and use lookup()
out_degrees=pld_edges_df.groupBy("src").count()
in_degrees=pld_edges_df.groupBy("dst").count()
pld_edges_df.unpersist()
out_degrees.show(3)
in_degrees.show(3)
#print(out_degrees.rdd.lookup(846558))
```

```
#print(fin_degrees_rdd.lookup(846558))
```

```
+---+-----+
|src|count|
```

```
+---+-----+
```

```
| 26|    2|
```

```
| 29|   34|
```

```
|964|    1|
```

```
+---+-----+
```

```
only showing top 3 rows
```

```
+-----+-----+
```

```
|      dst|count|
```

```
+-----+-----+
```

```
|      29|   40|
```

```
|36750820|    5|
```

```
|61427989| 3242|
```

```
+-----+-----+
```

```
only showing top 3 rows
```

```
%pyspark
```

ERROR

```
# Next, we'll construct a local dictionary from of all the PLDS (key is the PLD, value
# This is our truth-table of known PLDs that we'll use when counting hosts
# This code can't handle the full PLD list and produces this exception:
# Stack trace: ExitCodeException exitCode=52
#pld_lookup_table=dict(pld_df.rdd.map(lambda x: (x['PLD'], x['ID']))).collect())
#print(pld_lookup_table["aaa.aaa"])

# Instead, just create an RDD and use lookup()
#pld_lookup_table=pld_df.rdd.map(lambda x: (x['PLD'], x['ID']))
#print(pld_lookup_table.lookup("aaa.aaa"))

# Or let's try creating as a BloomFilter, since we only want to record presence of a PLI
expectedNumItems=91000000
fpp=0.005
#pld_bf = pld_df.stat.bloomFilter("PLD", expectedNumItems, fpp) # Doesn't exist in pyspark
#pld_bf.mightContain("aaa.aaa")
from pybloom import BloomFilter
pld_bf = BloomFilter(capacity=expectedNumItems, error_rate=fpp)
for row in pld_df.rdd.collect(): # TODO: Do this as a map with a function?
    pld_bf.add(row['PLD'])

print("aaa.aaa" in pld_bf)
print("aaa.aaa.bla" in pld_bf)

# Next, broadcast this map so it's available on all the slave nodes - this seems to bre
pld_bf_distrib=sc.broadcast(pld_bf)
```

```

File "<stdin>", line 5, in <module>
File "/usr/lib/spark/python/pyspark/rdd.py", line 809, in collect
    port = self.ctx._jvm.PythonRDD.collectAndServe(self._jrdd.rdd())
File "/usr/lib/spark/python/lib/py4j-0.10.4-src.zip/py4j/java_gateway.py", line 1133,
in __call__
    answer, self.gateway_client, self.target_id, self.name)
File "/usr/lib/spark/python/pyspark/sql/utils.py", line 63, in deco
    return f(*a, **kw)
File "/usr/lib/spark/python/lib/py4j-0.10.4-src.zip/py4j/protocol.py", line 319, in ge
t_return_value
    format(target_id, ".", name), value)
Py4JJavaError: An error occurred while calling z:org.apache.spark.api.python.PythonRDD.c
ollectAndServe.
: org.apache.spark.SparkException: Job aborted due to stage failure: Total size of seria
lized results of 70 tasks (8.2 GB) is bigger than spark.driver.maxResultSize (8.0 GB)
    at org.apache.spark.scheduler.DAGScheduler.org$apache$spark$scheduler$DAGSchedul
er$$failJobAndIndependentStages(DAGScheduler.scala:1690)
    at org.apache.spark.scheduler.DAGScheduler$$anonfun$abortStage$1.apply(DAGSchedu

```

```
%pyspark
```

ERROR

```

# Returns a Boolean to say whether PLD is a hostname in itself
def is_a_pld(hostname):
    #if hostname in pld_lookup_table:
    #if pld_lookup_table.filter(lambda a: a == hostname).count()>0:
    if hostname in pld_bf_distrib.value:
        return True
    else:
        return False

# Define a function to do the hostname->pld conversion, if the pld exists in our dictio
def convert_hostname(hostname):
    # Return hostname as-is, if this is already a PLD
    #if hostname in pld_lookup_table:
    #if pld_lookup_table.filter(lambda a: a == hostname).count()>0:
    if hostname in pld_bf_distrib.value:
        return hostname
    # Otherwise we're going to have to split it up and test the parts
    try:
        parts=hostname.split('.')
        if (len(parts)>4 and is_a_pld('.'.join(parts[0:4]))):
            return '.'.join(parts[0:4])
        if (len(parts)>3 and is_a_pld('.'.join(parts[0:3]))):
            return '.'.join(parts[0:3])
        if (len(parts)>2 and is_a_pld('.'.join(parts[0:2]))):
            return '.'.join(parts[0:2])
        if (len(parts)>1 and is_a_pld('.'.join(parts[0:1]))):
            return '.'.join(parts[0:1])
        return "ERROR" # Couldn't find a corresponding PLD - this should never happen!
    except:
        return "ERROR"

```

```
# Test
```

```
print(convert_hostname("aaa.aaa"))
```

Traceback (most recent call last):

```
File "/tmp/zeppelin_pyspark-7944601921887552320.py", line 367, in <module>
    raise Exception(traceback.format_exc())
```

Exception: Traceback (most recent call last):

```
File "/tmp/zeppelin_pyspark-7944601921887552320.py", line 355, in <module>
    exec(code, _zcUserQueryNameSpace)
```

```
File "<stdin>", line 22, in <module>
```

```
File "<stdin>", line 7, in convert_hostname
```

NameError: global name 'pld_bf_distrib' is not defined

```
%pyspark
```

READY

```
# Now count the number of hosts per PLD in a scalable way, and create another dictionary
# Takes 5mins for first 10M rows -> approx 8 hours for all 1.3B rows?
```

```
count_table=host_df.drop('hostid').rdd.map(lambda x: (convert_hostname(x['host']),1)).r
bool_table=host_df.drop('hostid').rdd.map(lambda x: (x['host'], is_a_pld(x['host']))).f
host_df.unpersist()
```

```
print(count_table['aaa.aaa'])
```

```
print(bool_table['aaa.aaa'])
```

```
print(count_table['ERROR']) # Should be zero once we've loaded all the PLDs!
```

```
# TODO: Fix error in collect()
```

```
# java.lang.IllegalArgumentException: Size exceeds Integer.MAX_VALUE
```

Traceback (most recent call last):

```
File "/tmp/zeppelin_pyspark-9191304803286420325.py", line 367, in <module>
    raise Exception(traceback.format_exc())
```

Exception: Traceback (most recent call last):

```
File "/tmp/zeppelin_pyspark-9191304803286420325.py", line 355, in <module>
    exec(code, _zcUserQueryNameSpace)
```

```
File "<stdin>", line 1, in <module>
```

```
File "/usr/lib/spark/python/pyspark/rdd.py", line 1570, in collectAsMap
    return dict(self.collect())
```

```
File "/usr/lib/spark/python/pyspark/rdd.py", line 809, in collect
```

```
port = self.ctx._jvm.PythonRDD.collectAndServe(self._jrdd.rdd())
```

```
File "/usr/lib/spark/python/lib/py4j-0.10.4-src.zip/py4j/java_gateway.py", line 1133,
in __call__
```

```
    answer, self.gateway_client, self.target_id, self.name)
```

```
File "/usr/lib/spark/python/pyspark/sql/utils.py", line 65, in deco
```

```
    s = e.java_exception.toString()
```

```
File "/usr/lib/spark/python/lib/py4j-0.10.4-src.zip/py4j/java_gateway.py", line 1131,
in _call
```

```
%pyspark
```

READY

```
from pyspark.sql.types import IntegerType
```

```
from pyspark.sql.functions import udf, col, when, lit
```

```
# Define a UDF to perform column-based lookup
def translate(mapping):
    def translate_(col):
        if not mapping.get(col):
            return 0
        else:
            return mapping.get(col)
    return udf(translate_, IntegerType())

# And a similar function for the Boolean map
def translate_bool(mapping):
    def translate_bool_(col):
        if not mapping.get(col):
            return False
        else:
            return mapping.get(col)
    return udf(translate_bool_, BooleanType())

# Insert our count column back into the host summary dataframe, along with a boolean to
# While we're at it, let's add in the in and out-degrees too, and an indicator of wheth
# crawled_test=when(col("OutDegree")==0, lit(False)).otherwise(lit(True))
pld_df_joined=pld_df.withColumn('NumHosts', translate(count_table)("PLD"))\
    .withColumn('PLDisHost?', translate_bool(bool_table)("PLD"))\
    #.withColumn('InDegree', translate(in_degrees)("ID"))\
    #.withColumn('OutDegree', translate(out_degrees)("ID"))\
    #.withColumn('Crawled?', crawled_test)

pld_df.unpersist()
pld_df_joined.sort("NumHosts", ascending=False).show(100)
```

Traceback (most recent call last):

```
File "/tmp/zeppelin_pyspark-9191304803286420325.py", line 367, in <module>
    raise Exception(traceback.format_exc())
```

Exception: Traceback (most recent call last):

```
File "/tmp/zeppelin_pyspark-9191304803286420325.py", line 355, in <module>
    exec(code, _zcUserQueryNameSpace)
```

```
File "<stdin>", line 17, in <module>
```

NameError: name 'count_table' is not defined

%pyspark

READY

```
# Join with in-degree and out-degree dataframes
pld_df_joined2=pld_df_joined.join(out_degrees, out_degrees.src==pld_df_joined.ID, "left")
pld_df_joined.unpersist()
pld_df_joined3=pld_df_joined2.join(in_degrees, in_degrees.dst==pld_df_joined2.ID, "left")
pld_df_joined2.unpersist()
pld_df_joined3.show(5)
pld_df_joined3.cache()
```

```
Traceback (most recent call last):
  File "/tmp/zeppelin_pyspark-9191304803286420325.py", line 367, in <module>
    raise Exception(traceback.format_exc())
Exception: Traceback (most recent call last):
  File "/tmp/zeppelin_pyspark-9191304803286420325.py", line 355, in <module>
    exec(code, _zcUserQueryNameSpace)
  File "<stdin>", line 1, in <module>
NameError: name 'pld_df_joined' is not defined
```

%pyspark

READY

```
# Insert a flag to indicate whether the PLD has been crawled
crawled_test=when(col("OutDegree").isNull(), lit(False)).otherwise(lit(True))
pld_df_joined4=pld_df_joined3.withColumn('Crawled?', crawled_test)
pld_df_joined3.unpersist()
pld_df_joined4.show(5)
pld_df_joined4.cache()
```

```
Traceback (most recent call last):
  File "/tmp/zeppelin_pyspark-9191304803286420325.py", line 367, in <module>
    raise Exception(traceback.format_exc())
Exception: Traceback (most recent call last):
  File "/tmp/zeppelin_pyspark-9191304803286420325.py", line 355, in <module>
    exec(code, _zcUserQueryNameSpace)
  File "<stdin>", line 2, in <module>
NameError: name 'pld_df_joined3' is not defined
```

%pyspark

READY

```
# Finally, join with the harmonic centrality and page-rank for each domain
# Note: could probably speed this up using something like above techniques, or by presorting
pld_df_joined5=pld_df_joined4.join(pr_df, pr_df.host_rev==pld_df_joined4.PLD, "leftOuter")
                                .withColumnRenamed("#hc_val", "HarmonicCentrality").withColumnRenamed("#pr_val", "PageRank")
pld_df_joined4.unpersist()
pld_df_joined5.show(5)
pld_df_joined5.cache()
```

```
Traceback (most recent call last):
  File "/tmp/zeppelin_pyspark-9191304803286420325.py", line 367, in <module>
    raise Exception(traceback.format_exc())
Exception: Traceback (most recent call last):
  File "/tmp/zeppelin_pyspark-9191304803286420325.py", line 355, in <module>
    exec(code, _zcUserQueryNameSpace)
  File "<stdin>", line 1, in <module>
NameError: name 'pld_df_joined4' is not defined
```

%pyspark

READY

```
# Save final table to S3 in compressed CSV format, broken into smaller files
```



```
outputURI="s3://billsdata.net/CommonCrawl/domain_summaries2/"  
codec="org.apache.hadoop.io.compress.GzipCodec"
```

Traceback (most recent call last):

```
File "/tmp/zeppelin_pyspark-9191304803286420325.py", line 367, in <module>  
    raise Exception(traceback.format_exc())
```

Exception: Traceback (most recent call last):

```
File "/tmp/zeppelin_pyspark-9191304803286420325.py", line 360, in <module>  
    exec(code, _zcUserQueryNameSpace)
```

```
File "<stdin>", line 3, in <module>
```

NameError: name 'pld_df_joined5' is not defined

%pyspark

READY