

Paul 5 - faster do...

%pyspark

FINISHED

```
# Zeppelin notebook to create domain summaries based on the May/Jun/Jul 2017 CommonCrawl
# as per description here: http://commoncrawl.org/2017/08/webgraph-2017-may-june-july/
# PJ - 7 October 2017
```

```
import boto
from pyspark.sql.types import *
```

```
LIMIT=1000000 # TODO - remove temporary limit to run full summaries!
```

```
# Import the PLD vertices list as a DataFrame
pld_schema=StructType([StructField("ID", StringType(), False), StructField("PLD", StringType(), False)])
pld_txt=sc.textFile("s3://commoncrawl/projects/hyperlinkgraph/cc-main-2017-may-jun-jul/pld-vertices.txt")
temp_pld = pld_txt.map(lambda k: k.split()) # By default, splits on whitespace, which is correct
pld_df=temp_pld.toDF(pld_schema).limit(LIMIT) #.repartition(4)
pld_df.show(3)
pld_df.cache()
# Should have 91M domains
#print(pld_df.count())
```

```
+---+-----+
| ID|    PLD|
+---+-----+
|  0|  aaa.a|
|  1|  aaa.aa|
|  2|aaa.aaa|
+---+-----+
only showing top 3 rows
DataFrame[ID: string, PLD: string]
```

%pyspark

FINISHED

```
# Next import the PLD edges as a DataFrame
pld_edges_schema=StructType([StructField("src", LongType(), False), StructField("dst", LongType(), False)])
pld_edges_txt=sc.textFile("s3://commoncrawl/projects/hyperlinkgraph/cc-main-2017-may-jun-jul/pld-edges.txt")
temp_edges_pld = pld_edges_txt.map(lambda k: map(int, k.split())) # By default, splits on whitespace, which is correct
pld_edges_df=temp_edges_pld.toDF(pld_edges_schema).limit(LIMIT*10) #.repartition(8)
pld_edges_df.show(3)
pld_edges_df.cache()
```

```

+---+-----+
|src|      dst|
+---+-----+
| 2| 9193244|
|20|75600973|
|21|46356172|
+---+-----+
only showing top 3 rows
DataFrame[src: bigint, dst: bigint]

```

%pyspark

FINISHED

```

# Load the host-level graph vertices in the same way
host_schema=StructType([StructField("hostid", StringType(), False), StructField("host",
host_txt=sc.textFile("s3://commoncrawl/projects/hyperlinkgraph/cc-main-2017-may-jun-jul.
temp_host = host_txt.map(lambda k: k.split()) # By default, splits on whitespace, which
host_df=temp_host.toDF(host_schema).limit(LIMIT*10).repartition(8)
host_df.show(3)
host_df.cache()
# Should have 1.3B hosts
#print(host_df.count())

```

```

+-----+-----+
|hostid|      host|
+-----+-----+
|    2|    aaa.aaa|
|   10|  aaa.aaa.www|
|   18|aaa.com.espaciola...|
+-----+-----+
only showing top 3 rows
DataFrame[hostid: string, host: string]

```

%pyspark

FINISHED

```

# Load in all harmonic centrality and page-ranks, and join based on reverse domain name
# Format: #hc_pos #hc_val #pr_pos #pr_val #host_rev
#pr_schema=StructType([StructField("hc_pos", StringType(), False), StructField("hc_val"
(), False), StructField("host_rev", StringType(), False)])
pr_txt=sc.textFile("s3://commoncrawl/projects/hyperlinkgraph/cc-main-2017-may-jun-jul/d
header=pr_txt.first()
pr_txt=pr_txt.filter(lambda x: x!=header)
temp_pr = pr_txt.map(lambda k: k.split()) # By default, splits on whitespace, which is
pr_df=temp_pr.toDF(header.split()).withColumnRenamed("#host_rev", "host_rev").limit(LIMI
pr_df.show(3)
pr_df.cache()
#print(pr_df.count()) # Should be 91M

```

```

+-----+-----+-----+-----+-----+
|#hc_pos| #hc_val|#pr_pos|          #pr_val|    host_rev|
+-----+-----+-----+-----+-----+
|      1|24989952|      1| 0.0155264576161686|  com.facebook|
|      2|22460880|      3|0.00866038900847366|  com.twitter|
|      3|22097514|      2| 0.0128827315785546|com.googleapis|
+-----+-----+-----+-----+-----+

```

only showing top 3 rows

```
DataFrame[#hc_pos: string, #hc_val: string, #pr_pos: string, #pr_val: string, host_rev: string]
```

```
%pyspark
```

FINISHED

```

# Debug partitioning of our 4 big dataframes
sc.getConf().getAll() #.mkString("\n")
print(pld_df.rdd.getNumPartitions())
print(pld_edges_df.rdd.getNumPartitions())
print(host_df.rdd.getNumPartitions())
pr_df.rdd.getNumPartitions()

```

```

1
1
8
1

```

```
%pyspark #--packages graphframes:graphframes:0.5.0-spark2.1-s_2.11
```

FINISHED

```
# We now have everything we need in these four dataframes to create the summaries we ne
```

```

# This code can't handle the complete edge lists, and produces this exception:
# java.lang.IllegalArgumentException: Size exceeds Integer.MAX_VALUE
#out_degrees=dict(pld_edges_df.groupBy("src").count().collect())
#in_degrees=dict(pld_edges_df.groupBy("dst").count().collect())
#print(out_degrees['846558'])
#print(in_degrees['846558'])

```

```

# Instead, just create RDDs and use lookup()
out_degrees=pld_edges_df.groupBy("src").count()
in_degrees=pld_edges_df.groupBy("dst").count()
pld_edges_df.unpersist()
out_degrees.show(3)
in_degrees.show(3)
#print(out_degrees.rdd.lookup(846558))
#print(in_degrees.rdd.lookup(846558))

```

```

+---+-----+
|src|count|
+---+-----+
| 2|    1|
|20|    1|
|21|    2|
+---+-----+
only showing top 3 rows
+-----+-----+
|    dst|count|
+-----+-----+
|9193244|   11|
|75600973|  21|
|46356172|  22|
+-----+-----+
only showing top 3 rows

```

```
%pyspark
```

FINISHED

```

# Next, we'll construct a local dictionary from of all the PLDS (key is the PLD, value
# This is our truth-table of known PLDs that we'll use when counting hosts
# This code can't handle the full PLD list and produces this exception:
# Stack trace: ExitCodeException exitCode=52
#pld_lookup_table=dict(pld_df.rdd.map(lambda x: (x['PLD'], x['ID']))).collect())
#print(pld_lookup_table["aaa.aaa"])

# Instead, just create an RDD and use lookup()
#pld_lookup_table=pld_df.rdd.map(lambda x: (x['PLD'], x['ID']))
#print(pld_lookup_table.lookup("aaa.aaa"))

# Or let's try creating as a BloomFilter, since we only want to record presence of a PLI
expectedNumItems=91000000
fpp=0.005
#pld_bf = pld_df.stat.bloomFilter("PLD", expectedNumItems, fpp) # Doesn't exist in pyspark
#pld_bf.mightContain("aaa.aaa")
from pybloom import BloomFilter
pld_bf = BloomFilter(capacity=expectedNumItems, error_rate=fpp)
for row in pld_df.rdd.collect():
    pld_bf.add(row['PLD'])

print("aaa.aaa" in pld_bf)
print("aaa.aaa.bla" in pld_bf)

# Next, broadcast this map so it's available on all the slave nodes - this seems to bre
pld_bf_distrib=sc.broadcast(pld_bf)

True
False

```

```
%pyspark
```

FINISHED

```
# Returns a Boolean to say whether PLD is a hostname in itself
```

```

def is_a_pld(hostname):
    #if hostname in pld_lookup_table:
    #if pld_lookup_table.filter(lambda a: a == hostname).count()>0:
    if hostname in pld_bf_distrib.value:
        return True
    else:
        return False

# Define a function to do the hostname->pld conversion, if the pld exists in our dictionary
def convert_hostname(hostname):
    # Return hostname as-is, if this is already a PLD
    #if hostname in pld_lookup_table:
    #if pld_lookup_table.filter(lambda a: a == hostname).count()>0:
    if hostname in pld_bf_distrib.value:
        return hostname
    # Otherwise we're going to have to split it up and test the parts
    try:
        parts=hostname.split('.')
        if (len(parts)>4 and is_a_pld('.'.join(parts[0:4]))):
            return '.'.join(parts[0:4])
        if (len(parts)>3 and is_a_pld('.'.join(parts[0:3]))):
            return '.'.join(parts[0:3])
        if (len(parts)>2 and is_a_pld('.'.join(parts[0:2]))):
            return '.'.join(parts[0:2])
        if (len(parts)>1 and is_a_pld('.'.join(parts[0:1]))):
            return '.'.join(parts[0:1])
        return "ERROR" # Couldn't find a corresponding PLD - this should never happen!
    except:
        return "ERROR"

# Test
print(convert_hostname("aaa.aaa"))
print(is_a_pld("aaa.aaa"))

aaa.aaa
True

```

```
%pyspark
```

ERROR

```

# Now count the number of hosts per PLD in a scalable way, and create another dictionary
# Takes 5mins for first 10M rows -> approx 8 hours for all 1.3B rows?
count_table=host_df.drop('hostid').rdd.map(lambda x: (convert_hostname(x['host']),1)).rdd.collect()
bool_table=host_df.drop('hostid').rdd.map(lambda x: (x['host'], is_a_pld(x['host']))).rdd.collect()
host_df.unpersist()
print(count_table['aaa.aaa'])
print(bool_table['aaa.aaa'])
print(count_table['ERROR']) # Should be zero once we've loaded all the PLDs!

# TODO: Fix error in collect()
# java.lang.IllegalArgumentException: Size exceeds Integer.MAX_VALUE

```

```

java:43)
  at java.lang.reflect.Method.invoke(Method.java:498)
  at py4j.reflection.MethodInvoker.invoke(MethodInvoker.java:244)
  at py4j.reflection.ReflectionEngine.invoke(ReflectionEngine.java:357)
  at py4j.Gateway.invoke(Gateway.java:280)
  at py4j.commands.AbstractCommand.invokeMethod(AbstractCommand.java:132)
  at py4j.commands.CallCommand.execute(CallCommand.java:79)
  at py4j.GatewayConnection.run(GatewayConnection.java:214)
  at java.lang.Thread.run(Thread.java:748)
Caused by: org.apache.spark.api.python.PythonException: Traceback (most recent call last
):
  File "/mnt/yarn/usercache/zeppelin/appcache/application_1507360501747_0005/container_1
507360501747_0005_01_000003/pyspark.zip/pyspark/worker.py", line 177, in main
    process()
  File "/mnt/yarn/usercache/zeppelin/appcache/application_1507360501747_0005/container_1
507360501747_0005_01_000003/pyspark.zip/pyspark/worker.py", line 172, in process
    serializer.dump_stream(func(split_index, iterator), outfile)
  File "/usr/lib/spark/python/pyspark/rdd.py", line 2423, in pipeline_func

```

```
%pyspark
```

ERROR

```

from pyspark.sql.types import IntegerType
from pyspark.sql.functions import udf, col, when, lit

# Define a UDF to perform column-based lookup
def translate(mapping):
    def translate_(col):
        if not mapping.get(col):
            return 0
        else:
            return mapping.get(col)
    return udf(translate_, IntegerType())

# And a similar function for the Boolean map
def translate_bool(mapping):
    def translate_bool_(col):
        if not mapping.get(col):
            return False
        else:
            return mapping.get(col)
    return udf(translate_bool_, BooleanType())

# Insert our count column back into the host summary dataframe, along with a boolean to
# While we're at it, let's add in the in and out-degrees too, and an indicator of wheth
# crawled_test=when(col("OutDegree")==0, lit(False)).otherwise(lit(True))
pld_df_joined=pld_df.withColumn('NumHosts', translate(count_table)("PLD"))\
    .withColumn('PLDisHost?', translate_bool(bool_table)("PLD"))\
    #.withColumn('InDegree', translate(in_degrees)("ID"))\
    #.withColumn('OutDegree', translate(out_degrees)("ID"))\
    #.withColumn('Crawled?', crawled_test)

pld_df.unpersist()
pld_df_joined.sort("NumHosts", ascending=False).show(100)
pld_df_joined.cache()

```

Traceback (most recent call last):

```
File "/tmp/zeppelin_pyspark-7103816659305902723.py", line 367, in <module>
    raise Exception(traceback.format_exc())
```

Exception: Traceback (most recent call last):

```
File "/tmp/zeppelin_pyspark-7103816659305902723.py", line 355, in <module>
    exec(code, _zcUserQueryNameSpace)
```

```
File "<stdin>", line 17, in <module>
```

NameError: name 'count_table' is not defined

%pyspark

ERROR

```
# Join with in-degree and out-degree dataframes
pld_df_joined2=pld_df_joined.join(out_degrees, out_degrees.src==pld_df_joined.ID, "left")
pld_df_joined.unpersist()
pld_df_joined3=pld_df_joined2.join(in_degrees, in_degrees.dst==pld_df_joined2.ID, "left")
pld_df_joined2.unpersist()
pld_df_joined3.show(5)
pld_df_joined3.cache()
```

Traceback (most recent call last):

```
File "/tmp/zeppelin_pyspark-7103816659305902723.py", line 367, in <module>
    raise Exception(traceback.format_exc())
```

Exception: Traceback (most recent call last):

```
File "/tmp/zeppelin_pyspark-7103816659305902723.py", line 355, in <module>
    exec(code, _zcUserQueryNameSpace)
```

```
File "<stdin>", line 1, in <module>
```

NameError: name 'pld_df_joined' is not defined

%pyspark

ERROR

```
# Insert a flag to indicate whether the PLD has been crawled
crawled_test=when(col("OutDegree").isNull(), lit(False)).otherwise(lit(True))
pld_df_joined4=pld_df_joined3.withColumn('Crawled?', crawled_test)
pld_df_joined3.unpersist()
pld_df_joined4.show(5)
pld_df_joined4.cache()
```

Traceback (most recent call last):

```
File "/tmp/zeppelin_pyspark-7103816659305902723.py", line 367, in <module>
    raise Exception(traceback.format_exc())
```

Exception: Traceback (most recent call last):

```
File "/tmp/zeppelin_pyspark-7103816659305902723.py", line 355, in <module>
    exec(code, _zcUserQueryNameSpace)
```

```
File "<stdin>", line 2, in <module>
```

NameError: name 'pld_df_joined3' is not defined

%pyspark

ERROR

```
# Finally, join with the harmonic centrality and page-rank for each domain
# Note: could probably speed this up using something like above techniques, or by presorting
pld_df_joined5=pld_df_joined4.join(pr_df, pr_df.host_rev==pld_df_joined4.PLD, "leftOuter")
                                .withColumnRenamed("#hc_val", "HarmonicCentrality").withColumnRenamed("#pr_val", "PageRank")
pld_df_joined4.unpersist()
pld_df_joined5.show(5)
pld_df_joined5.cache()
```

Traceback (most recent call last):

```
File "/tmp/zeppelin_pyspark-7103816659305902723.py", line 367, in <module>
    raise Exception(traceback.format_exc())
```

Exception: Traceback (most recent call last):

```
File "/tmp/zeppelin_pyspark-7103816659305902723.py", line 355, in <module>
    exec(code, _zcUserQueryNamespace)
```

```
File "<stdin>", line 1, in <module>
```

NameError: name 'pld_df_joined4' is not defined

%pyspark

ERROR

```
# Save final table to S3 in compressed CSV format
outputURI="s3://billsdata.net/CommonCrawl/domain_summaries2/"
codec="org.apache.hadoop.io.compress.GzipCodec"
pld_df_joined5.coalesce(1).write.format('com.databricks.spark.csv').options(header='true', mode='overwrite')
```

Traceback (most recent call last):

```
File "/tmp/zeppelin_pyspark-7103816659305902723.py", line 367, in <module>
    raise Exception(traceback.format_exc())
```

Exception: Traceback (most recent call last):

```
File "/tmp/zeppelin_pyspark-7103816659305902723.py", line 360, in <module>
    exec(code, _zcUserQueryNamespace)
```

```
File "<stdin>", line 3, in <module>
```

NameError: name 'pld_df_joined5' is not defined

%pyspark

FINISHED