

# Tom 2 Wiki Topic ...

%pyspark

FINISHED

```
# PySpark CommonCrawl Topic Modelling
# Tom V / Paul J - 13/2/2018
```

```
# SET THE spark.driver.maxResultSize PROPERTY TO 3G
```

```
import boto
from boto.s3.key import Key
from gzipstream import GzipStreamFile
from pyspark.sql.types import *
import warc
import ujson as json
from urlparse import urlparse
from langdetect import detect_langs
import pyclld2 as cld2
```

```
#wetlist = sc.textFile("s3://commoncrawl/crawl-data/CC-MAIN-2017-04/wet.paths.gz") # Apr
# Latest blog/documentation: http://commoncrawl.org/2017/10/october-2017-crawl-archive-1
wetlist = sc.textFile("s3://commoncrawl/crawl-data/CC-MAIN-2017-43/wet.paths.gz") # Oct
```

```
wetlist.cache()
```

```
def unpack(uri):
    conn = boto.connect_s3(anon=True, host='s3.amazonaws.com')
    bucket = conn.get_bucket('commoncrawl')
    key_ = Key(bucket, uri)
    file_ = warc.WARCFile(fileobj=GzipStreamFile(key_))
    return file_
```

```
def detect(x):
    try:
        return detect_langs(x)[0].lang # Maybe we can get away with looking at less cha
    except Exception as e:
        return None
```

```
def detect2(x):
    try:
        isReliable, textBytesFound, details = cld2.detect(x)
        return details[0][1]
    except Exception as e:
        print(e)
        return None
```

```
def process_wet(id_, iterator):
    for uri in iterator:
        file = unpack(uri)
        for record in file: # Approx 53k web pages per WET file
            try:
                #url = record.rec_headers.get_header('WARC-Target-URI')
```

```

        #yield record, record.content_stream().read().decode('utf-8')
        url = record.url

        # TODO: Limit number of bytes read per record e.g. read(200000)

        domain = None if not url else urlparse(url).netloc
        text = record.payload.read().decode('utf-8') #.limit(100) # TODO: Limit
        lang = detect2(text[:300]) # Use PyCLD2, not langdetect, which was kill
        yield domain, url, text, lang
    except Exception as e:
        yield e

def process_wet_simple(id_, iterator):
    count=0
    for uri in iterator:
        file = unpack(uri)
        for record in file:
            try:
                count=count+1
                # TODO: Output total size of pages, rather than number of pages
                # Histogram.
            except Exception as e:
                pass
        #print(count)

```

```

%pyspark
detect2("this is a test")

'en'

```

READY

```

%pyspark

# PARAMETER - number of input files
nfiles = 1 # Total 89100

# PARAMETER - slices / partitions of input
files = sc.parallelize(wetlist.take(nfiles)) #, numSlices=nfiles/32) # TODO: Try numSli

# Should parallelize
print(files.getNumPartitions())
rdd=files.mapPartitionsWithIndex(process_wet)

print(str(rdd))
docs = rdd.toDF(["host", "url", "text","lang"]) # "lang"
#docs.cache()
#docs.count() # Total docs in all languages

320
PythonRDD[102] at RDD at PythonRDD.scala:48

```

READY

```

%pyspark

```

READY

```
# Filter for English only
docs_en = docs.filter(docs.lang == 'en')
```

## Load saved vectors from Wikipedia model (created by python Wikipedia Text Processing.ipynb)

READY

```
%pyspark
from pyspark.ml import Pipeline, PipelineModel
from pyspark.ml.feature import RegexTokenizer, CountVectorizer, StopWordsRemover
from pyspark.ml.clustering import LocalLDAModel

textModel = PipelineModel.load('s3://billsdata.net/CommonCrawl/wikipedia/text_model')
ldaModel = LocalLDAModel.load('s3://billsdata.net/CommonCrawl/wikipedia/lda_model')
```

```
%pyspark
```

READY

```
# Test the models - for debugging only
import numpy as np
import pandas as pd

X=ldaModel.topicsMatrix().toArray()
vocab = np.array(textModel.stages[2].vocabulary)

topicLabels = [' '.join(vocab[np.argsort(X[:,i])[:, :-1]][:5])] for i in range(100)]

def score_topics(text):
    df = sqlContext.createDataFrame(pd.DataFrame({'text':[text]}))
    vec = textModel.transform(df)
    scores = ldaModel.transform(vec).select('topicDistribution').collect()[0].topicDist
    return pd.Series(dict(zip(topicLabels, scores)))

# Try it on an arbitrary sentence
print(score_topics("This is the latest news about North Korea and their involvement in
school students education university college
season team first teams cup
series book published books novel
series show television also episode
ship ships two navy war
social one may also people
space earth light solar star
species found also large may
station line railway service train
team season coach football first
tom oliver ghost haiti kay
ukrainian ukraine dog dogs stamps
university research professor published science
war union soviet communist political
water company construction new coal
world olympics championships summer women
zealand new grand auckland prix
Length: 100, dtype: float64")
```

`%pyspark`

READY

```
# Now score pages from our WET files
docs_en.show(5)
vec = textModel.transform(docs_en)
vec.show(5)
```

```
+-----+-----+-----+-----+
|          null|          null|Software-Info: ia...| en|
|1000daysofwriting...|http://1000daysof...|1000 Days of Writ...| en|
|100unhappydays.bl...|http://100unhappy...|100 Unhappy Days:...| en|
|          10in30.com|http://10in30.com...|LearnOutLoud_300x...| en|
|123-free-download...|http://123-free-d...|MusicBoxTool - [3...| en|
+-----+-----+-----+-----+
```

only showing top 5 rows

```
+-----+-----+-----+-----+
-+-----+-----+-----+-----+
|          host|          url|          text|lang|          word
s|          filtered|          vec|
+-----+-----+-----+-----+
-+-----+-----+-----+-----+
|          null|          null|Software-Info: ia...| en|[software, info, ..
.|[software, info, ...|(20000,[88,152,33...|
|1000daysofwriting...|http://1000daysof...|1000 Days of Writ...| en|[days, of, writin..
.|[days, writing, d...|(20000,[0,2,3,5,7...|
```

`%pyspark`

READY

```
# Create topic distribution vectors and tidy up
scores = ldaModel.transform(vec)
scores2 = scores.drop("url").drop("text").drop("lang").drop("words").drop("filtered").d
scores2.show(5)
```

```
+-----+-----+-----+
|          host| topicDistribution|
+-----+-----+-----+
|          null|[0.13351995547840...|
|1000daysofwriting...|[3.26238961502545...|
|100unhappydays.bl...|[5.81230792144732...|
|          10in30.com|[8.42785330446217...|
|123-free-download...|[6.44166735802645...|
+-----+-----+-----+
```

only showing top 5 rows

`%pyspark`

READY

```
# Save these vectors to disc, so we can just load them later
scores2.write.parquet('s3://billsdata.net/CommonCrawl/topic_model_%d_files/cc_page_wiki.
```

**Load saved scores from nfiles of WET files**

FINISHED

`%pyspark``# Restart here``nfiles=1``scores2 = spark.read.parquet('s3://billsdata.net/CommonCrawl/topic_model_%d_files/cc_pa  
scores2.show(5)`

```

+-----+-----+
|          domain|  topicDistribution|
+-----+-----+
|          null|[0.12615256587902...|
|1000daysofwriting...|[3.26238961502545...|
|100unhappydays.bl...|[5.81230792144733...|
|          10in30.com|[8.42785330446217...|
|123-free-download...|[6.44166735802645...|
+-----+-----+
only showing top 5 rows

```

`%pyspark`

FINISHED

`# Aggregate page-scores per Host for now (will be same process for aggregating host-sco  
scores3=scores2.rdd.map(lambda x: (x['domain'], (1,x['topicDistribution']))).reduceByKey``# Next, divide by the total to create averaged vectors, and convert back to a dataframe  
scores4=scores3.map(lambda x: (x[0], (x[1][1]/x[1][0]))).toDF(["host", "averageTopicDis  
scores4.show(5)``{s.name: type(s.dataType) for s in scores4.schema}`

```

+-----+-----+
|          host|averageTopicDistribution|
+-----+-----+
| www.superdrug.com| [0.01613454968703...|
|      kiteforum.com| [1.11136619419678...|
|www.virtualmuseum.ca| [1.80597394573963...|
|ukrainianstartups...| [2.46670116225146...|
| lauragravestock.com| [4.49484962755617...|
+-----+-----+
only showing top 5 rows
{'averageTopicDistribution': <class 'pyspark.ml.linalg.VectorUDT'>, 'host': <class 'pysp  
ark.sql.types.StringType'>}

```

`%pyspark`

FINISHED

`# Just playing - code to help understand the different libraries and vector types!``import pyspark.mllib.linalg as mllib``import pyspark.ml.linalg as ml``df = sc.parallelize([  
 (mllib.DenseVector([1, ]), ml.DenseVector([1, ])),`

```

    (mllib.SparseVector(1, [0, ], [1, ]), ml.SparseVector(1, [0, ], [1, ]))
  ]).toDF(["mllib_v", "ml_v"])
df.show()
{s.name: type(s.dataType) for s in df.schema}

+-----+-----+
|      mllib_v|      ml_v|
+-----+-----+
|      [1.0]|      [1.0]|
|(1,[0],[1.0])|(1,[0],[1.0])|
+-----+-----+
{'ml_v': <class 'pyspark.ml.linalg.VectorUDT'>, 'mllib_v': <class 'pyspark.mllib.linalg.VectorUDT'>}
```

%pyspark

FINISHED

```

# TODO: Enrich each row with the corresponding PLD (using code from Paul J, but pickle
saveURI="s3://billsdata.net/CommonCrawl/hyperlinkgraph/cc-main-2017-aug-sep-oct/domaing
pld_df=spark.read.load(saveURI)
pld_df.show(3)
pld_df.cache()
```

```

+---+-----+
| ID|  PLD|
+---+-----+
| 0|aaa.1|
| 1|aaa.2|
| 2|aaa.3|
+---+-----+
only showing top 3 rows
DataFrame[ID: string, PLD: string]
```

%pyspark

ERROR

```

# Next, we'll construct a local dictionary from of all the PLDS (key is the PLD, value
# This is our truth-table of known PLDs that we'll use when counting hosts
# Create a bloom filter using a pure python package (might be a little slow)
from pybloom import BloomFilter
pld_bf = BloomFilter(capacity=94000000, error_rate=0.005) # was 91M

for row in pld_df.rdd.collect(): # limit(10000000) # TODO: Still bad (and exceeds spark
    pld_bf.add(row['PLD'])

print(pld_df.rdd.take(3))
print(pld_df.rdd.take(3)[2]['PLD'])
#pld_bf.add(pld_df.rdd.take(3)[2]['PLD'])
print("aaa.aaa" in pld_bf) # Should be True

import sys
print(sys.getsizeof(pld_bf))
print(len(pld_bf)) # Should match number of items entered
```

```

# Broadcast the bloom filter so it's available on all the slave nodes - we don't need to
# it any more so it's fine being immutable.
pld_bf_distrib=sc.broadcast(pld_bf)

print("aaa.aaa" in pld_bf) # Should be true
print("aaa.aaa.bla" in pld_bf) # Should be false
print("aaa.aaa" in pld_bf_distrib.value) # Should be true
print("aaa.aaa.bla" in pld_bf_distrib.value) # Should be false

org.apache.thrift.transport.TTransportException
  at org.apache.thrift.transport.TIOStreamTransport.read(TIOStreamTransport.java:1
32)
    at org.apache.thrift.transport.TTransport.readAll(TTransport.java:86)
    at org.apache.thrift.protocol.TBinaryProtocol.readAll(TBinaryProtocol.java:429)
    at org.apache.thrift.protocol.TBinaryProtocol.readI32(TBinaryProtocol.java:318)
    at org.apache.thrift.protocol.TBinaryProtocol.readMessageBegin(TBinaryProtocol.j
ava:219)
    at org.apache.thrift.TServiceClient.receiveBase(TServiceClient.java:69)
    at org.apache.zepplin.interpreter.thrift.RemoteInterpreterService$Client.recv_i
nterpret(RemoteInterpreterService.java:266)
    at org.apache.zepplin.interpreter.thrift.RemoteInterpreterService$Client.interp
ret(RemoteInterpreterService.java:250)
    at org.apache.zepplin.interpreter.remote.RemoteInterpreter.interpret(RemoteInte
rpreter.java:373)
    at org.apache.zepplin.interpreter.LazyOpenInterpreter.interpret(LazyOpenInterpr
eter.java:97)
    at org.apache.zepplin.notebook.Paragraph.jobRun(Paragraph.java:406)
    at org.apache.zepplin.scheduler.Job.run(Job.java:175)
    at org.apache.zepplin.scheduler.RemoteScheduler$JobRunner.run(RemoteScheduler.j
ava:329)
    at java.util.concurrent.Executors$RunnableAdapter.call(Executors.java:511)
    at java.util.concurrent.FutureTask.run(FutureTask.java:266)
    at java.util.concurrent.ScheduledThreadPoolExecutor$ScheduledFutureTask.access$2
01(ScheduledThreadPoolExecutor.java:180)
    at java.util.concurrent.ScheduledThreadPoolExecutor$ScheduledFutureTask.run(Sche
duledThreadPoolExecutor.java:293)
    at java.util.concurrent.ThreadPoolExecutor.runWorker(ThreadPoolExecutor.java:114
9)
    at java.util.concurrent.ThreadPoolExecutor$Worker.run(ThreadPoolExecutor.java:62
4)
    at java.lang.Thread.run(Thread.java:748)

```

```
%pyspark
```

ERROR

```

from pyspark.sql.functions import udf

# Returns a Boolean to say whether PLD is a hostname in itself
def is_a_pld(hostname):
    #if hostname in pld_lookup_table:
    #if pld_lookup_table.filter(lambda a: a == hostname).count()>0:
    if hostname in pld_bf_distrib.value:
        return True
    else:

```

```

    return False

# Define a function to do the hostname->pld conversion, if the pld exists in our dictio
def convert_hostname(hostname):
    # Return hostname as-is, if this is already a PLD
    #if hostname in pld_lookup_table:
    #if pld_lookup_table.filter(lambda a: a == hostname).count()>0:
    if hostname in pld_bf_distrib.value:
        return hostname
    # Otherwise we're going to have to split it up and test the parts
    try:
        parts=hostname.split('.')
        if (len(parts)>4 and is_a_pld('.'.join(parts[0:4]))):
            return '.'.join(parts[0:4])
        if (len(parts)>3 and is_a_pld('.'.join(parts[0:3]))):
            return '.'.join(parts[0:3])
        if (len(parts)>2 and is_a_pld('.'.join(parts[0:2]))):
            return '.'.join(parts[0:2])
        if (len(parts)>1 and is_a_pld('.'.join(parts[0:1]))):
            return '.'.join(parts[0:1])
        return "ERROR" # Couldn't find a corresponding PLD - this should never happen!
    except:
        return "ERROR"

udf_convert_hostname = udf(convert_hostname, StringType())

# Test
print(convert_hostname("aaa.aaa"))
#print(is a pld("aaa.aaa")) # Should be true

org.apache.thrift.transport.TTransportException
    at org.apache.thrift.transport.TIOStreamTransport.read(TIOStreamTransport.java:1
32)
    at org.apache.thrift.transport.TTransport.readAll(TTransport.java:86)
    at org.apache.thrift.protocol.TBinaryProtocol.readAll(TBinaryProtocol.java:429)
    at org.apache.thrift.protocol.TBinaryProtocol.readI32(TBinaryProtocol.java:318)
    at org.apache.thrift.protocol.TBinaryProtocol.readMessageBegin(TBinaryProtocol.j
ava:219)
    at org.apache.thrift.TServiceClient.receiveBase(TServiceClient.java:69)
    at org.apache.zepplin.interpreter.thrift.RemoteInterpreterService$Client.recv_i
nterpret(RemoteInterpreterService.java:266)
    at org.apache.zepplin.interpreter.thrift.RemoteInterpreterService$Client.interp
ret(RemoteInterpreterService.java:250)
    at org.apache.zepplin.interpreter.remote.RemoteInterpreter.interpret(RemoteInte
rpreter.java:373)
    at org.apache.zepplin.interpreter.LazyOpenInterpreter.interpret(LazyOpenInterpr
eter.java:97)
    at org.apache.zepplin.notebook.Paragraph.jobRun(Paragraph.java:406)
    at org.apache.zepplin.scheduler.Job.run(Job.java:175)
    at org.apache.zepplin.scheduler.RemoteScheduler$JobRunner.run(RemoteScheduler.j
ava:329)
    at java.util.concurrent.Executors$RunnableAdapter.call(Executors.java:511)
    at java.util.concurrent.FutureTask.run(FutureTask.java:266)
    at java.util.concurrent.ScheduledThreadPoolExecutor$ScheduledFutureTask.access$2
01(ScheduledThreadPoolExecutor.java:180)
    at java.util.concurrent.ScheduledThreadPoolExecutor$ScheduledFutureTask.run(Sche

```



```

duledThreadPoolExecutor.java:293)
    at java.util.concurrent.ThreadPoolExecutor.runWorker(ThreadPoolExecutor.java:114
9)
    at java.util.concurrent.ThreadPoolExecutor$Worker.run(ThreadPoolExecutor.java:62
4)
    at java.lang.Thread.run(Thread.java:748)

```

```
%pyspark
```

ERROR

```

# Function to reverse hostnames
from pyspark.sql.functions import udf
def reverse_domain(domain):
    return '.'.join(reversed(domain.split('.')))
print(reverse_domain("com.facebook"))
udf_reverse_domain = udf(reverse_domain, StringType())

# Convert hosts in Topic DF to PLDs using convert_hostname function from Paul 5.
scores5=scores4.withColumn("host_rev",udf_reverse_domain(udf_convert_hostname(udf_rever:
scores5.show(10)

```

```

java.net.SocketException: Broken pipe (Write failed)
    at java.net.SocketOutputStream.socketWrite0(Native Method)
    at java.net.SocketOutputStream.socketWrite(SocketOutputStream.java:111)
    at java.net.SocketOutputStream.write(SocketOutputStream.java:155)
    at java.io.BufferedOutputStream.flushBuffer(BufferedOutputStream.java:82)
    at java.io.BufferedOutputStream.write(BufferedOutputStream.java:121)
    at org.apache.thrift.transport.TIOStreamTransport.write(TIOStreamTransport.java:
145)
    at org.apache.thrift.protocol.TBinaryProtocol.writeString(TBinaryProtocol.java:2
02)
    at org.apache.zepplin.interpreter.thrift.RemoteInterpreterContext$RemoteInterpr
eterContextStandardScheme.write(RemoteInterpreterContext.java:1133)
    at org.apache.zepplin.interpreter.thrift.RemoteInterpreterContext$RemoteInterpr
eterContextStandardScheme.write(RemoteInterpreterContext.java:992)
    at org.apache.zepplin.interpreter.thrift.RemoteInterpreterContext.write(RemoteI
nterpreterContext.java:882)
    at org.apache.zepplin.interpreter.thrift.RemoteInterpreterService$interpret_arg
s$interpret_argsStandardScheme.write(RemoteInterpreterService.java:6501)
    at org.apache.zepplin.interpreter.thrift.RemoteInterpreterService$interpret_arg
s$interpret_argsStandardScheme.write(RemoteInterpreterService.java:6424)
    at org.apache.zepplin.interpreter.thrift.RemoteInterpreterService$interpret_arg
s.write(RemoteInterpreterService.java:6351)
    at org.apache.thrift.TServiceClient.sendBase(TServiceClient.java:63)
    at org.apache.zepplin.interpreter.thrift.RemoteInterpreterService$Client.send_i
nterpret(RemoteInterpreterService.java:260)
    at org.apache.zepplin.interpreter.thrift.RemoteInterpreterService$Client.interp
ret(RemoteInterpreterService.java:249)
    at org.apache.zepplin.interpreter.remote.RemoteInterpreter.interpret(RemoteInte
rpreter.java:373)
    at org.apache.zepplin.interpreter.LazyOpenInterpreter.interpret(LazyOpenInterpr
eter.java:97)
    at org.apache.zepplin.notebook.Paragraph.jobRun(Paragraph.java:406)
    at org.apache.zepplin.scheduler.Job.run(Job.java:175)

```

```

    at org.apache.zeppelin.scheduler.RemoteScheduler$JobRunner.run(RemoteScheduler.java:329)
    at java.util.concurrent.Executors$RunnableAdapter.call(Executors.java:511)
    at java.util.concurrent.FutureTask.run(FutureTask.java:266)
    at java.util.concurrent.ScheduledThreadPoolExecutor$ScheduledFutureTask.access$201(ScheduledThreadPoolExecutor.java:180)
    at java.util.concurrent.ScheduledThreadPoolExecutor$ScheduledFutureTask.run(ScheduledThreadPoolExecutor.java:293)
    at java.util.concurrent.ThreadPoolExecutor.runWorker(ThreadPoolExecutor.java:1149)
    at java.util.concurrent.ThreadPoolExecutor$Worker.run(ThreadPoolExecutor.java:624)
    at java.lang.Thread.run(Thread.java:748)

```

%pyspark

ERROR

```

# TODO: Now we can aggregate page-scores per PLD, using a map-reduce similar to the hos
java.net.ConnectException: Connection refused (Connection refused)
    at java.net.PlainSocketImpl.socketConnect(Native Method)
    at java.net.AbstractPlainSocketImpl.doConnect(AbstractPlainSocketImpl.java:350)
    at java.net.AbstractPlainSocketImpl.connectToAddress(AbstractPlainSocketImpl.java:206)
    at java.net.AbstractPlainSocketImpl.connect(AbstractPlainSocketImpl.java:188)
    at java.net.SocksSocketImpl.connect(SocksSocketImpl.java:392)
    at java.net.Socket.connect(Socket.java:589)
    at org.apache.thrift.transport.TSocket.open(TSocket.java:182)
    at org.apache.zeppelin.interpreter.remote.ClientFactory.create(ClientFactory.java:51)
    at org.apache.zeppelin.interpreter.remote.ClientFactory.create(ClientFactory.java:37)
    at org.apache.commons.pool2.BasePooledObjectFactory.makeObject(BasePooledObjectFactory.java:60)
    at org.apache.commons.pool2.impl.GenericObjectPool.create(GenericObjectPool.java:861)
    at org.apache.commons.pool2.impl.GenericObjectPool.borrowObject(GenericObjectPool.java:435)
    at org.apache.commons.pool2.impl.GenericObjectPool.borrowObject(GenericObjectPool.java:363)
    at org.apache.zeppelin.interpreter.remote.RemoteInterpreterProcess.getClient(RemoteInterpreterProcess.java:92)
    at org.apache.zeppelin.interpreter.remote.RemoteInterpreter.interpret(RemoteInterpreter.java:352)
    at org.apache.zeppelin.interpreter.LazyOpenInterpreter.interpret(LazyOpenInterpreter.java:97)
    at org.apache.zeppelin.notebook.Paragraph.jobRun(Paragraph.java:406)
    at org.apache.zeppelin.scheduler.Job.run(Job.java:175)
    at org.apache.zeppelin.scheduler.RemoteScheduler$JobRunner.run(RemoteScheduler.java:329)
    at java.util.concurrent.Executors$RunnableAdapter.call(Executors.java:511)
    at java.util.concurrent.FutureTask.run(FutureTask.java:266)
    at java.util.concurrent.ScheduledThreadPoolExecutor$ScheduledFutureTask.access$2

```

```

01(ScheduledThreadPoolExecutor.java:180)
    at java.util.concurrent.ScheduledThreadPoolExecutor$ScheduledFutureTask.run(Sche
    duledThreadPoolExecutor.java:293)
    at java.util.concurrent.ThreadPoolExecutor.runWorker(ThreadPoolExecutor.java:114
9)
    at java.util.concurrent.ThreadPoolExecutor$Worker.run(ThreadPoolExecutor.java:62
4)
    at java.lang.Thread.run(Thread.java:748)

```

```
%pyspark
```

READY

```

# TODO: Save pld topic distributions in parquet format for Tom to play with (and to fig
# Maybe a numpy argmax to get the index of the 'top' topic for each PLD with a score.

```